


Qwen2.5-Math-7B · Training reward

Bold lines: trailing 20-step mean · Faint lines: per-step reward



Reward: correct = 1; incorrect = -0.5; unextracted = -1. One seed per arm.

Microbatch: Baseline 32; Pro-only / Max-only 16. Snapshot: 2026-09-26 19:08:46 UTC+8